

LLM Evaluation Metrics

Sergey Tarasenko, PhD

ROUGE

(Recall-Oriented Understudy for Gisting Evaluation)

The Core Concept: Precision, Recall, and F1

At its heart, ROUGE measures overlap using basic evaluation principles:

- **Recall:** How much of the human summary did the AI manage to capture?

$$\text{Recall} = \frac{\text{Matching Words}}{\text{Total Words in Human Reference Summary}}$$

- **Precision:** How much of the AI's summary was actually relevant/correct?

$$\text{Precision} = \frac{\text{Matching Words}}{\text{Total Words in AI Summary}}$$

- **F1-Score:** The harmonic mean that balances both precision and recall.

ROUGE (Recall-Oriented Understudy for Gisting Evaluation) is a set of metrics used in Natural Language Processing to evaluate automatically generated text—most commonly **summaries** and **translations**.

It works by comparing an AI-generated summary (candidate) against a human-written summary (reference) to see how much overlapping content they share.

Formally, ROUGE-N is an n-gram recall between a candidate summary and a set of reference summaries. ROUGE-N is computed as follows:

$$\text{ROUGE-N} = \frac{\sum_{S \in \{\text{ReferenceSummaries}\}} \sum_{gram_n \in S} \text{Count}_{match}(gram_n)}{\sum_{S \in \{\text{ReferenceSummaries}\}} \sum_{gram_n \in S} \text{Count}(gram_n)} \quad (1)$$

Worked Example

- **Human Reference:** *"The cat sat on the mat."* (6 words)
- **AI Summary:** *"The cat sat."* (3 words)

Calculating ROUGE-1:

- **Matching words:** "The", "cat", "sat" (3 matches)
- **Recall:** $3/6 = 0.50$ (50%)
- **Precision:** $3/3 = 1.00$ (100%)
- **F1-Score:** $2 \times \frac{0.5 \times 1.0}{0.5 + 1.0} = 0.67$ (67%)

BLUE

How BLEU Works

BLEU calculates a score from **0.0 to 1.0** (or 0 to 100%) based on two main components:

1. Modified n-gram Precision

BLEU looks at individual words (unigrams), two-word sequences (bigrams), three-word sequences (trigrams), and four-word sequences (4-grams) to see what percentage of the AI's generated phrases appear in the human reference.

To prevent a model from cheating by simply repeating a correct word over and over (e.g., repeating "*the the the*"), BLEU **clips** word counts so a word can only match up to the maximum number of times it appears in the reference.

$$\text{Modified Precision}_n = \frac{\sum \text{Clipped Count of Matching } n\text{-grams}}{\sum \text{Total } n\text{-grams in AI Text}}$$

2. Brevity Penalty (BP)

Because precision naturally favors shorter outputs (a 1-word translation that happens to be correct would get a 100% precision score), BLEU penalizes outputs that are shorter than the human reference text:

$$\text{BP} = \begin{cases} 1 & \text{if candidate length } c > \text{reference length } r \\ e^{(1-r/c)} & \text{if candidate length } c \leq \text{reference length } r \end{cases}$$

BLEU vs. ROUGE Comparison

Feature	BLEU	ROUGE
Primary Focus	Precision (Are generated words accurate?)	Recall (Did we capture all ground-truth words?)
Primary Domain	Machine Translation, Code Generation	Summarization, QA
N-gram Target	Averages 1-gram through 4-gram precision	Often measured individually (ROUGE-1, ROUGE-2, ROUGE-L)
Length Control	Uses a Brevity Penalty for short text	Naturally favors longer texts unless using F1-score

Limitations with Modern Large Language Models (LLMs)

While BLEU remains popular for classic machine translation benchmarks, it has notable limitations for modern LLM evaluation:

- **Exact Match Dependency:** Like ROUGE, BLEU cannot recognize synonyms or semantically identical paraphrases.
- **Brittle with Open-Ended Prompts:** For creative writing, chat responses, or general reasoning, there is rarely a single "correct" reference text to match against.
- **Modern Alternatives:** LLM evaluation increasingly uses semantic embeddings (BERTScore) or model-based judges (LLM-as-a-Judge like GPT-4 or Claude to grade responses based on reasoning quality).

GLUE

(General Language Understanding Evaluation)

Unlike ROUGE and BLEU—which are single statistical formulas measuring word overlap—**GLUE** (General Language Understanding Evaluation) is a comprehensive **benchmark suite** of 9 distinct tasks designed to evaluate a model's general natural language understanding (NLU).

Instead of scoring raw text generation (like summarization or translation), GLUE tests how well a model grasps grammar, sentiment, sentence similarity, and logical relationships.

The 9 Tasks in GLUE

GLUE aggregates a model's performance across three core natural language understanding categories into a single composite score (0 to 100):

1. Single-Sentence Tasks

- **CoLA** (Corpus of Linguistic Acceptability): Tests whether a sentence is grammatically correct.
- **SST-2** (Stanford Sentiment Treebank): Binary sentiment classification (predicting positive vs. negative movie reviews).

2. Similarity and Paraphrase Tasks

- **MRPC** (Microsoft Research Paraphrase Corpus): Determines if two sentences are semantically equivalent.
- **QQP** (Quora Question Pairs): Identifies if two questions asked on Quora have the same meaning.
- **STS-B** (Semantic Textual Similarity Benchmark): Scores the semantic similarity of two sentence pairs on a scale from 1 to 5.

3. Inference and Logic Tasks (NLI)

- **MNLI** (Multi-Genre Natural Language Inference): Given a premise and hypothesis, determines if the hypothesis is *entailed* (true), *contradictory* (false), or *neutral*.
- **QNLI** (Question Natural Language Inference): Evaluates whether a given context sentence contains the answer to a question.
- **RTE** (Recognizing Textual Entailment): A smaller textual entailment dataset covering various news and Wikipedia sources.
- **WNLI** (Winograd NLI): Tests pronoun resolution (e.g., figuring out what "it" refers to in a complex sentence).

SuperGLUE

SuperGLUE: The Harder Successor

When GLUE was released in 2018, models like BERT and RoBERTa rapidly surpassed estimated human performance. To prevent benchmark saturation, researchers created **SuperGLUE** in 2019, introducing harder tasks requiring deeper reasoning:

- **Coreference Resolution** (WSC)
- **Question Answering in Context** (MultiRC, ReCoRD)
- **Causal Reasoning** (COPA)
- **Complex Entailment** (CB, RTE)

How GLUE Relates to Modern Generative LLMs

Feature	GLUE / SuperGLUE	Modern LLM Benchmarks (MMLU, GSM8K)
Model Type	Designed for encoder models (BERT, RoBERTa)	Designed for generative decoder models (GPT-4, Llama)
Evaluation Style	Fine-tuning model heads for classification	Zero-shot or few-shot prompt generation
Primary Scope	Core linguistic understanding and grammar	Multi-domain knowledge, coding, and multi-step reasoning

BIG-bench (Beyond Imitation Game)

BIG-bench (short for **Beyond the Imitation Game Benchmark**) is one of the largest, most comprehensive collaborative evaluation suites created to test and probe the limits of Large Language Models (LLMs).

Initiated by Google researchers in 2021, the goal of BIG-bench was to move past basic language tests and build a benchmark diverse enough to uncover **emergent capabilities**—behaviors, reasoning skills, and pitfalls that only appear when models reach massive scale.

It consists of over **204 distinct tasks** testing a wide array of capabilities:

- **Complex Reasoning:** Multi-step arithmetic, symbolic logic, temporal logic, and spatial reasoning.
- **Niche Knowledge & Linguistics:** Translation of low-resource or synthetic languages, deciphering riddles, and analyzing jargon.
- **Creative & Social Dynamics:** Movie identification from emojis, recognizing sarcasm, theory of mind, and understanding metaphors.
- **Safety & Ethics:** Probing for gender/racial bias, calibration (does the model know when it doesn't know?), and social stereotypes.

MMLU
(Massive Multitask Language Understanding)

MMLU consists of **15,908 multiple-choice questions** spanning **57 diverse subjects**.

The test is usually evaluated in a **few-shot** (typically 5-shot) or **zero-shot** setting. In a 5-shot evaluation, the LLM is given 5 solved example questions from that subject as context, followed by the test question, and must output the correct choice (A, B, C, or D).

Core Subject Categories

Category	Examples of Included Subjects
STEM	Abstract Algebra, College Physics, Virology, Computer Security
Humanities	Philosophy, World Religions, Prehistory, High School History
Social Sciences	Econometrics, Jurisprudence, US Foreign Policy, Sociology
Professional	Professional Medicine, Clinical Knowledge, Professional Law, Accounting

HELM

(Holistic Evaluation of Language Models)

HELM (Holistic Evaluation of Language Models) is not a single mathematical formula—it is a comprehensive, multi-dimensional evaluation framework and benchmark created by Stanford University's Center for Research on Foundation Models (CRFM).

HELM was developed because traditional benchmarks (like MMLU, GLUE, or ROUGE) offer a narrow view of performance. Chasing a single metric like "accuracy" can obscure critical real-world failure modes like toxicity, bias, or poor calibration.

HELM evaluates models using a two-dimensional grid: Scenarios crossed with Metrics.

HELM tests models across 42+ scenarios representing real-world use cases, including:

- Core NLP Tasks: Summarization (CNN/DailyMail), Information Extraction, Sentiment Analysis.
- Reasoning & QA: Mathematics (GSM8K, MATH), Code (HumanEval), Legal Reasoning (LSAT), Medical QA.
- Aspirational Scenarios: Multi-turn dialogue, non-native English text, regional dialects, and copyrighted text detection.

The 7 Primary Metrics Measured in HELM

Metric	What It Evaluates	Why It Matters
Accuracy	Direct correctness on tasks (using reference answers, exact matches, etc.).	Core model capability.
Calibration	Alignment between the model's self-reported confidence and actual correctness.	Ensures the model knows when it is uncertain instead of hallucinating confidently.
Robustness	Performance stability when inputs are slightly altered (e.g., typos, rephrasing, prompt variation).	Tests real-world durability against messy user inputs.
Fairness	Consistency in performance across demographic groups or dialect variations.	Prevents performance degradation for underrepresented users.
Bias	Measurement of stereotypical associations or asymmetric output behavior.	Prevents the propagation of harmful real-world biases.
Toxicity	Rate of generating harmful, offensive, or hateful output (e.g., using RealToxicityPrompts).	Crucial for user safety and deployment eligibility.
Efficiency	Computational cost, memory usage, and token latency.	Determines production costs and deployment viability.

Language Model Metrics Evolution

○ BLEU (Bilingual Evaluation Understudy)

2002

Introduced by IBM researchers (Papineni et al., 2002), **BLEU** measures **precision** by calculating modified n -gram overlap between generated text and human reference texts. Originally designed for Machine Translation, it became an early standard across natural language generation tasks.

○ ROUGE (Recall-Oriented Understudy for Gisting Evaluation)

2004

Created by Chin-Yew Lin (2004), **ROUGE** shifted the focus to **recall** by measuring how much of the human reference content was captured in the model's output. It quickly became the primary metric for evaluating text summarization.

○ GLUE (General Language Understanding Evaluation)

2018

Developed by researchers at NYU, UW, and DeepMind (Wang et al., 2018), **GLUE** aggregated 9 diverse Natural Language Understanding (NLU) tasks (sentiment analysis, textual entailment, question answering) into a single unified benchmark suite. It shifted evaluation from single-task metrics to measuring general task transfer.

○ SuperGLUE

2019

Released as a direct successor (Wang et al., 2019) after transformer models rapidly surpassed human baselines on GLUE. **SuperGLUE** introduced significantly harder tasks, including complex coreference resolution, multi-sentence reasoning, and casual reasoning with more demanding scoring criteria.

○ MMLU (Massive Multitask Language Understanding)

2020

Introduced by Dan Hendrycks et al. (2020), **MMLU** covers 57 subjects across STEM, humanities, social sciences, and professional domains using 4-option multiple-choice questions. It became the primary benchmark for measuring broad world knowledge and zero/few-shot reasoning in modern LLMs.

○ BIG-bench (Beyond the Imitation Game Benchmark)

2022

Initiated by Google Research (Srivastava et al., 2022 / 2023) involving 450+ authors across 132 institutions. **BIG-bench** comprised 204+ diverse, complex tasks probing emergent behaviors, specialized capabilities, logic, and limitations in large language models.



HELM (Holistic Evaluation of Language Models)

2022

Created by Percy Liang et al. (Bommasani et al., 2023) at Stanford's Center for Research on Foundation Models (CRFM), **HELM** moved beyond pure accuracy. It evaluates models across **7 core metrics** (accuracy, calibration, robustness, fairness, bias, toxicity, and efficiency) across 40+ scenarios to provide a comprehensive profile of LLM capabilities and risks.

Key Paradigm Shifts

- **2002–2004 (Lexical Matching):** Focused on surface-level n -gram overlaps (precision/recall) against human ground truth.
- **2018–2019 (Multi-Task NLU):** Shifted toward evaluating deep linguistic comprehension across diverse NLP datasets.
- **2020–2022 (Massive Probing & Zero-Shot Knowledge):** Expanded to multi-thousand task suites and multi-subject testing to capture emergent capabilities and broad reasoning.
- **2022+ (Holistic & Multi-Dimensional):** Expanded to safety, toxicity, calibration, and robustness alongside functional capability.

References

- Bommasani, R., Liang, P., & Lee, T. (2023). Holistic Evaluation of Language Models. *Annals of the New York Academy of Sciences*, 1525(1), 140–146. <https://doi.org/10.1111/nyas.15007>
- Hendrycks, D., Burns, C., Basart, S., Zou, A., Mazeika, M., Song, D., & Steinhardt, J. (2020). Measuring Massive Multitask Language Understanding. *arXiv preprint arXiv:2009.03300*.
- Lin, C.-Y. (2004). ROUGE: A Package for Automatic Evaluation of Summaries. In *Text Summarization Branches Out: Proceedings of the ACL-04 Workshop* (pp. 74–81). Association for Computational Linguistics.
- Papineni, K., Roukos, S., Ward, T., & Zhu, W.-J. (2002). BLEU: a Method for Automatic Evaluation of Machine Translation. In *Proceedings of the 40th Annual Meeting on Association for Computational Linguistics* (pp. 311–318). Association for Computational Linguistics. <https://doi.org/10.3115/1073083.1073135>
- Srivastava, A., Rastogi, A., Rao, A., Shoeb, A. A. M., Abid, A., Fisch, A., ... & Wu, Z. (2023). Beyond the Imitation Game: Quantifying and extrapolating the capabilities of language models. *Transactions on Machine Learning Research (TMLR)*. <https://openreview.net/forum?id=uyTL5Bvosj>
- Wang, A., Singh, A., Michael, J., Hill, F., Levy, O., & Bowman, S. (2018). GLUE: A Multi-Task Benchmark and Analysis Platform for Natural Language Understanding. In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP* (pp. 353–355). Association for Computational Linguistics. <https://doi.org/10.18653/v1/w18-5446>
- Wang, A., Pruksachatkun, Y., Nangia, N., Singh, A., Michael, J., Hill, F., Levy, O., & Bowman, S. R. (2019). SuperGLUE: A Stickier Benchmark for General-Purpose Language Understanding Systems. *arXiv preprint arXiv:1905.00537*.

